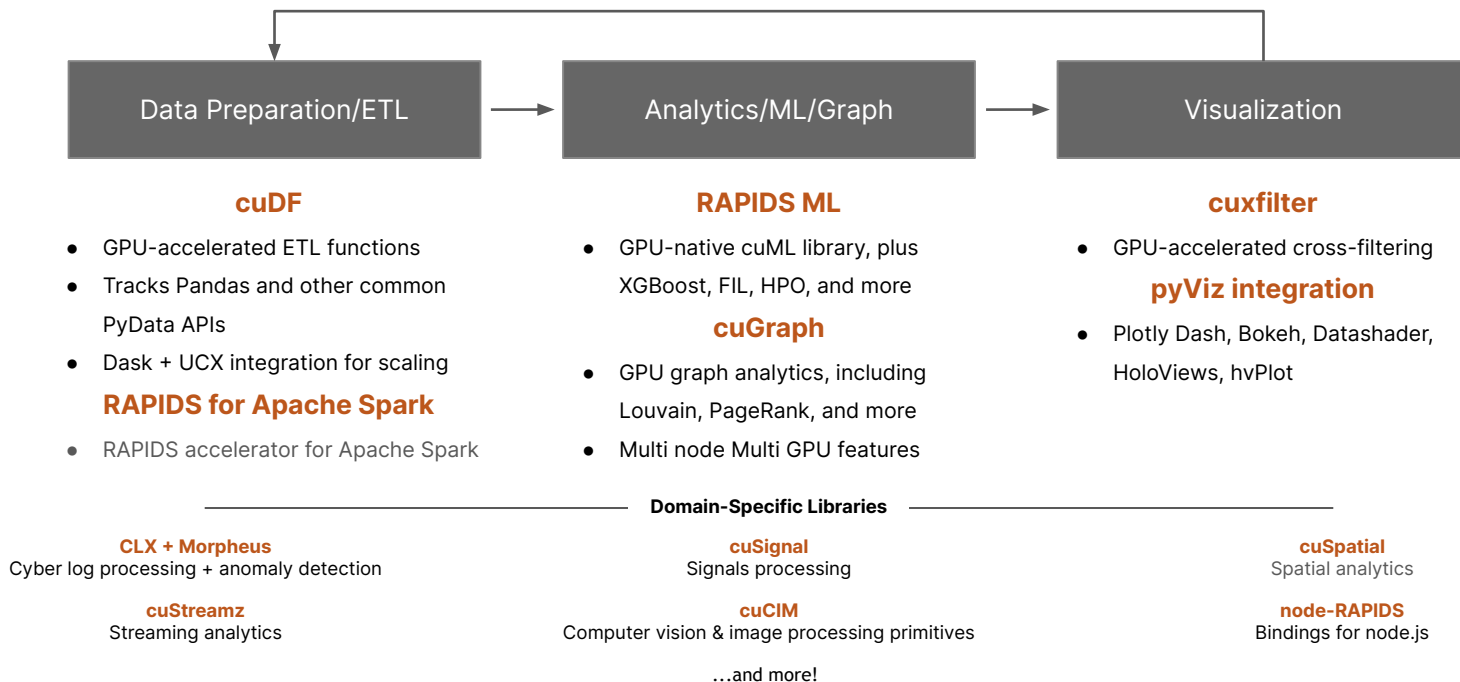


RAPIDS

The RAPIDS suite of open source software libraries gives you the freedom to execute end-to-end data science and analytics pipelines entirely on NVIDIA GPUs.


RAPIDS Product Ecosystem




RAPIDS Growth

More than **1.5M** downloads and accelerating

RAPIDS Downloads

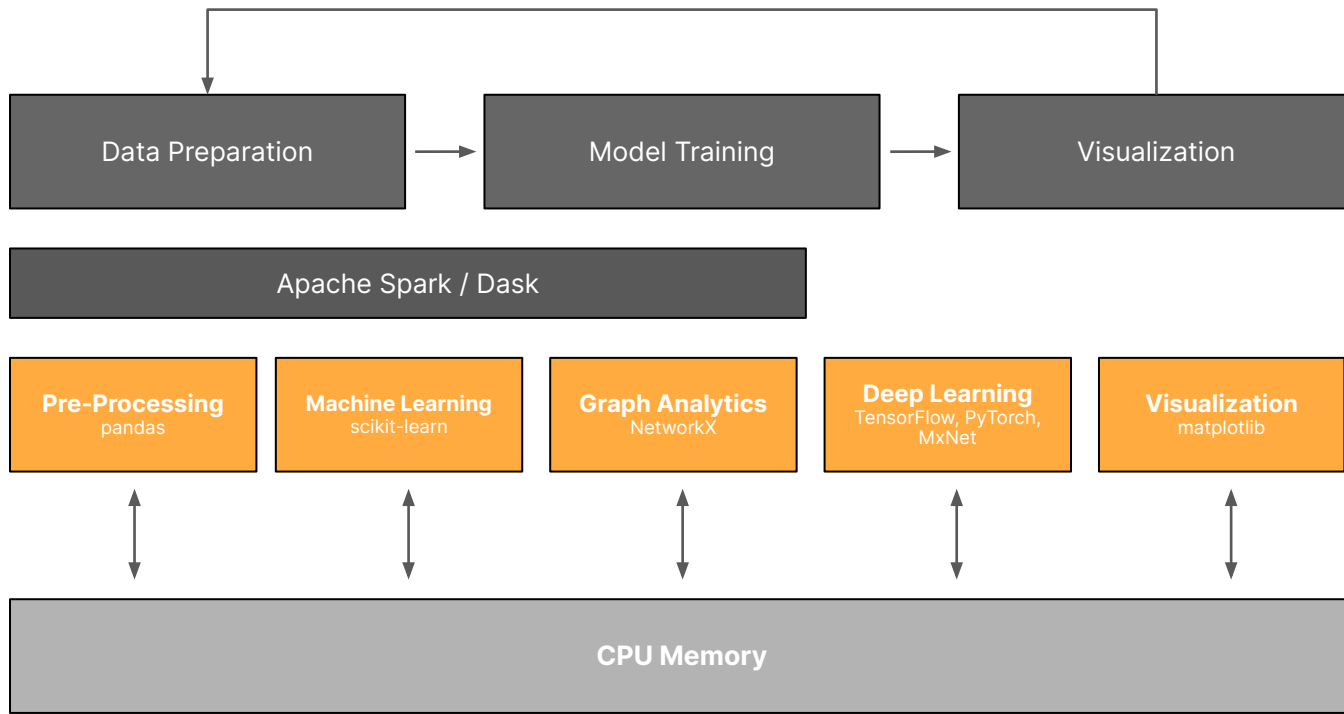


RAPIDS Community Engagement



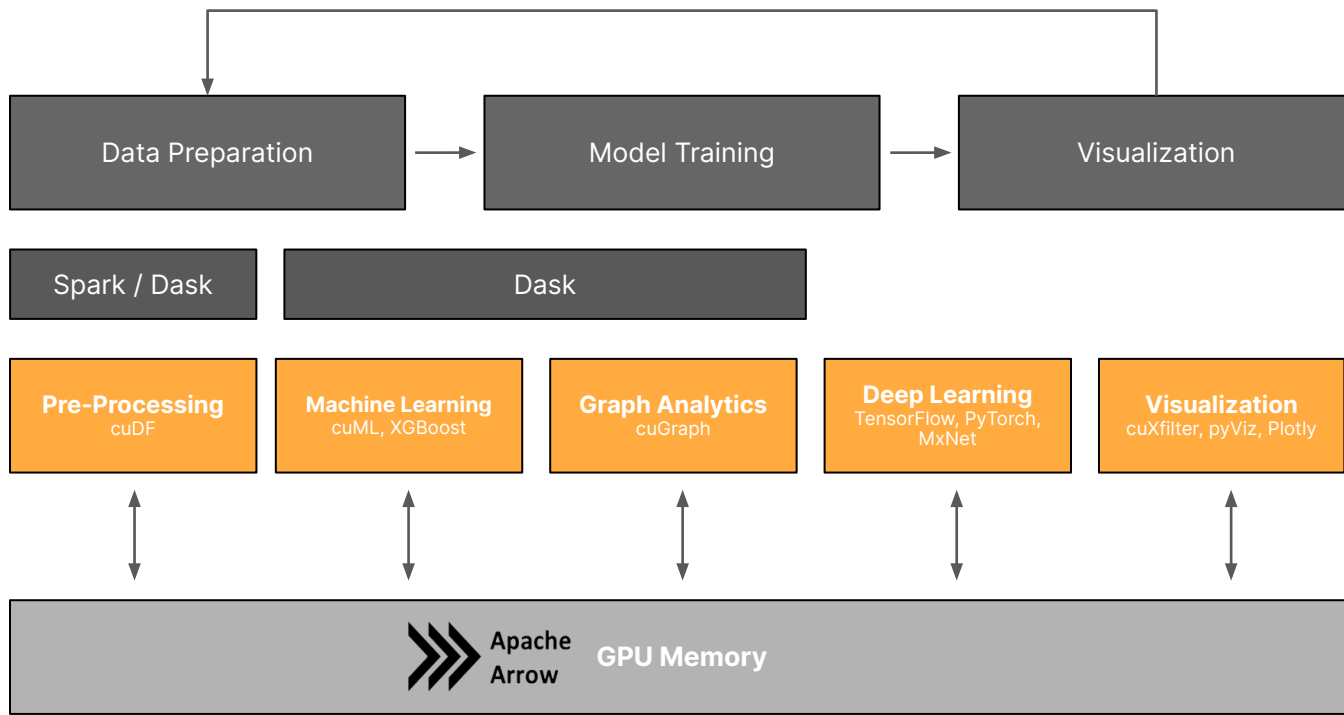
Open Source Software Has Democratized Data Science

Highly Accessible, Easy to Use Tools Abstract Complexity



Accelerated Data Science with RAPIDS

Powering Popular Data Science Ecosystems with NVIDIA GPUs



Minor Code Changes for Major Benefits

Abstracting Accelerated Compute through Familiar Interfaces

CPU

pandas

```
>>> import pandas as pd
>>> df =
pd.read_csv("filepath")
```

CPU Spark

```
spark.sql("""
select
    order
    count(*) as
order_count
from
    orders""")
```

scikit-learn

```
>>> from sklearn.ensemble
import
RandomForestClassifier
>>> clf =
RandomForestClassifier()
>>> clf.fit(x, y)
```

NetworkX

```
>>> import networkx as nx
>>> page_rank =
nx.pagerank(graph)
```



GPU

cuDF

```
>>> import cudf
>>> df =
cudf.read_csv("filepath")
```

Average Speed-Ups: 150x

GPU Spark

```
spark.conf.set("spark.plugins
", "com.nvidia.spark.SQLPlugin")
spark.sql("""
select
    order
    count(*) as order_count
from
    orders""")
```

Average Speed-Ups: 10x

cuML

```
>>> from cuml.ensemble
import
RandomForestClassifier
>>> cuclf =
RandomForestClassifier()
>>> cuclf.fit(x, y)
```

Average Speed-Ups: 50x

cuGraph

```
>>> import cugraph
>>> page_rank =
cugraph.pagerank(graph)
```

Average Speed-Ups: 250x

The Evolution of Data Processing

Faster Data Access, Less Data Movement

Hadoop Processing, Reading from Disk




CPU-Based Spark In-Memory Processing




25-100x Improvement
Less Code
Language Flexible
Primarily In-Memory

Traditional GPU Processing



5-10x Improvement
More Code
Language Rigid
Substantially on GPU

RAPIDS



50-100x Improvement
Same Code
Language Flexible
Primarily on GPU

?


Why Use GPUs

GPUs are built for intensive parallel processing. As datasets continue to grow, data scientists are limited by the sequential nature of CPU compute. GPUs provide the power and parallelism necessary for today's data science.

Lightning-Fast End-to-End Performance

Reducing Data Science Processes from Hours to Seconds

RAPIDS End-to-End Workflow Runtimes



16

A100s Provide More Power than 100 CPU Nodes

70x

Faster Performance than Similar CPU Configuration

20x

More Cost-Effective than Similar CPU Configuration

Demo

Using RAPIDS on Hyperplane